

Geometry of Multi-Turn Conversation Trajectories in Small Language Models: Characterization, Signal-Conditioned Compression, and Limits

Pranav Kompally

Independent Research Checkpoint

New York, USA

`rt-geometry-memory repository checkpoint`

Abstract—We present a manuscript-grade checkpoint of a three-stage research program on conversational memory compression for small language models, together with an explicit record of where the program succeeds and where it fails. The program proceeds from geometry measurement (Paper 1) to geometry-aware control under fixed memory budgets (Paper 2) to sparse signal-conditioned conversational memory codecs (Paper 3). In Paper 1, conversation-state trajectories are extremely low-rank across three compact model families (mean rank95 1.049–1.528) and geometric distortion is tightly coupled to decoder drift (geometry-logit correlation 0.989–0.994), with permutation controls collapsing both effects. In Paper 2, geometry-aware budget allocation beats uniform retention on a custom hard stress set by rescuing memory-critical support turns more often (17/36 vs. 5/36 on `qwen25_05b` at budget 0.35). In Paper 3, we establish that the keep/compress/drop (KCD) sparse codec is a signal-agnostic compression structure whose optimal scoring depends on the memory type the task requires. Four decisive experiments close this question. First, a direct signal comparison on the hard stress set shows that `geometry_KCD` significantly outperforms `semantic_KCD` on constraint-critical memory (pairwise $\Delta\Delta_{\ell_2}$ of +37.8 at budget 0.20, $p = 0.028$; support-turn rescue 17/36 vs. 7/36), confirming that the geometry signal is load-bearing and cannot be replaced by semantic scoring when exact constraint packets must be preserved. Second, a Gate 1 oracle probe on MSC valid (16 conversations, 23,928 candidate rows) shows that query-conditioned geometric features add substantial ranking value inside a semantic shortlist (Δ Kendall tau +0.53–0.60, far exceeding the +0.03 decision threshold), and that `semantic_query_conditioned_geometry_KCD` beats the strongest semantic-only codec at budget 0.35 on both logit distortion ($p = 0.044$ conversation) and answer NLL ($p = 0.020$). Third, the same Gate 1 probe on LongMemEval-S cleaned (6 conversations, 10,416 candidate rows) passes at 17–23 $\times$ the threshold (Δ Kendall +0.51–0.53), confirming the cross-benchmark generality of the geometry ranking advantage. Fourth, a 3B model validation shows the geometry advantage grows with scale: on `qwen25_3b`, `semantic_KCD` produces +1.13 nats of answer NLL damage at budget 0.50 ($p = 0.0003$) relative to `geometry_KCD`, while every geometry policy improves answer quality relative to uniform. We also report two negative results from further investigated open items—a learned harm predictor and memory-object level compression—that define precise future work directions, and five earlier falsification results that define the claim boundary.

Index Terms—large language models, long-context memory,

conversational memory, representation geometry, memory compression, sparse codecs, signal-conditioned compression

I. INTRODUCTION

Long multi-turn conversations create a memory problem before they create a generation problem. Once the full dialogue no longer fits inside a practical inference budget, the model must decide what to retain, what to compress, and what to evict. The central claim of this repository checkpoint is that *hidden-state geometry of the evolving conversation trajectory* is a useful and sometimes necessary control signal for that decision—but the right way to use it depends on the type of memory the task requires.

This repository organizes the project into three linked papers:

- 1) **Paper 1: Geometry characterization.** Determine whether turn-level conversation-state trajectories are low-rank, curved in informative ways, and predictively tied to decoder behavior.
- 2) **Paper 2: Geometry-aware control.** Use geometry-derived risk scores to allocate a fixed memory budget.
- 3) **Paper 3: Signal-conditioned sparse codecs.** Replace part of raw history retention with a compressed memory object whose scoring signal is matched to the memory type of the task.

The current checkpoint completes all three papers. Paper 1 is frozen; Paper 2 has a stable mechanism story; Paper 3 now has decisive evidence on both sides of the geometry-versus-semantics gap—geometry is necessary and sufficient for constraint-critical memory, and geometry adds significant value inside semantic shortlists for persona-continuity memory. The main contribution is no longer “geometry beats baselines.” It is that the KCD codec framework with signal-dependent scoring is the right architectural response to the diversity of conversational memory types.

a) *Main contributions of the checkpoint.:*

- A geometric formalization of turn-level conversational dynamics on the sphere using local tangent-space transport.

Checkpoint Research Program Overview




Fig. 1. Research program overview. Paper 1 establishes low-rank, decoder-relevant conversation-state geometry. Paper 2 converts geometry into a budget allocation signal. Paper 3 turns that signal into a signal-conditioned sparse keep/compress/drop memory codec whose scoring adapts to memory type.

- Stable empirical evidence that conversation-state trajectories are extremely low-rank and strongly decoder-relevant across compact model families.
- A geometry-aware memory-control result under conversational scarcity, including a mechanism analysis showing rescue of memory-critical support turns.
- A signal-conditioned sparse codec family (KCD) that is active, budget-responsive, and mechanism-valid, with decisive evidence that the signal is load-bearing.
- A completed Gate 1 oracle experiment on both MSC valid (Δ Kendall +0.53–0.60) and LongMemEval-S cleaned (Δ Kendall +0.51–0.53), confirming geometry adds substantial ranking value inside semantic shortlists across benchmark families.
- A 3B model-scale validation showing geometry advantage grows with representational capacity: `semantic_KCD` at 3B is catastrophically harmful (+1.13 NLL, $p = 0.0003$) while geometry-based policies consistently improve answer quality.
- Two investigated open items yielding informative negatives: a learned harm predictor (trains well but cross-conversation miscalibration prevents it from beating geometry at policy evaluation) and memory-object compression (mechanical under-retention failure; diagnoses a clear fix direction).
- A falsification record with five failed directions that sharpen the claim boundary: failed persona-vs.-filler curvature separation on MSC, a too-coarse geometry-only regime atlas, and three failed state-update supersession formulations.

II. MATHEMATICAL SETUP

A. Conversation-State Geometry

For a multi-turn conversation $x_{1:T}$, let

$$z_t = f_\theta(x_{1:t}) \in \mathbb{R}^d \quad (1)$$

be a turn-level hidden-state summary extracted from the model after turn t . We normalize onto the sphere:

$$h_t = \frac{z_t}{\|z_t\|} \in \mathbb{S}^{d-1}. \quad (2)$$

Within a short segment S_j , choose a reference point $q_j \in \mathbb{S}^{d-1}$. The one-step spherical increment from h_t to h_{t+1} is mapped into the tangent space at q_j :

$$u_t = \text{PT}_{h_t \rightarrow q_j}(\log_{h_t}(h_{t+1})) \in T_{q_j} \mathbb{S}^{d-1}, \quad (3)$$

where $\log$ is the Riemannian logarithm map and PT is parallel transport. Paper 1 tests whether these transported increments are approximately low-rank:

$$u_t \approx B_j c_t + e_t, \quad B_j \in \mathbb{R}^{d \times r_j}, \quad r_j \ll d. \quad (4)$$

We use the following discrete curvature proxy, stabilized with an arclength floor to avoid blow-up on near-stationary segments:

$$\kappa_t \approx \frac{\|u_t - u_{t-1}\|_2}{\max(\Delta s_t, \epsilon)}, \quad (5)$$

with Δs_t the local geodesic step length and ϵ a small floor. A large κ_t indicates locally unstable geometry or a potential conversational regime transition.

B. Query-Conditioned Geometry

For Paper 3’s semantic-memory extension, we project the geometry into the query’s local tangent frame. Given a target turn q with state h_q , define the query-projected curvature and local subspace energy:

$$\kappa_t^{(q)} = \kappa_t \cdot \langle \hat{u}_t, \text{PT}_{h_q \rightarrow q_j}(\hat{u}_q) \rangle, \quad (6)$$

where $\hat{u}$ denotes unit-normalized increments. The query-conditioned risk score $r_t^{(q)}$ combines query-projected curvature with query-projected local subspace energy in the same tangent frame as the conversation motion. This makes geometry query-aware without abandoning the spherical transport frame.

C. Decoder-Relevance Metrics

Given a reconstructed state trajectory $\hat{h}_{1:T}$, decoder distortion is measured by comparing the full and compressed decoder outputs:

$$\Delta_{\ell_2} = \|\ell(\hat{h}_t) - \ell(h_t)\|_2, \quad (7)$$

$$\text{KL} = \text{KL}\left(p_\theta(\cdot | h_t) \parallel p_\theta(\cdot | \hat{h}_t)\right), \quad (8)$$

$$\text{Top1} = \mathbf{1} \left[\arg \max p_\theta(\cdot | h_t) = \arg \max p_\theta(\cdot | \hat{h}_t) \right]. \quad (9)$$

Paper 2 and Paper 3 additionally evaluate answer-level average negative log-likelihood relative to the gold next assistant response.

D. Budgeted Memory Control

For a memory policy π acting on conversation units u , the control problem is

$$\min_{\pi} \mathbb{E}[D(\pi)] \quad \text{s.t.} \quad \mathbb{E}[M(\pi)] \leq B, \quad (10)$$

where $D(\pi)$ is decoder distortion and $M(\pi)$ is realized memory cost. Paper 2 instantiates π as a turn- or segment-level retention rule driven by a geometry score. Paper 3 turns π into a sparse memory-action policy

$$a_j \in \{\text{keep, compress, evict}\} \quad (11)$$

over segment j , producing a compressed memory object

$$m_j = (a_j, s_j, \text{meta}_j), \quad (12)$$

where a_j is the anchor/keyframe summary, s_j is sparse support memory, and meta_j stores role/order information. The scoring signal feeding the KCD selector is a free parameter: geometric risk for constraint-critical tasks, semantic shortlist with query-conditioned geometric refinement for persona-continuity tasks.

III. THEORY DIRECTIONS AND PROOF SKETCHES

A. Paper 1: Transported Low-Rank Geometry

Proposition 1. *Let $h_t \in \mathbb{S}^{d-1}$ be a conversation-state trajectory and let u_t be its transported increments in a common tangent space $T_{q_j}\mathbb{S}^{d-1}$. Then local low-rank analysis of the sequence $\{u_t\}_{t \in S_j}$ is well-defined independently of the original ambient coordinates, up to the choice of segment reference point and the local transport gauge.*

Proof sketch. Parallel transport is an isometry between tangent spaces on the sphere. Therefore norms and principal angles between transported increments are preserved under the transport map. Once all increments in a segment are moved into the same tangent space, ordinary Euclidean low-rank analysis becomes meaningful without mixing incompatible local coordinates. $\square$

Theorem 1. *Assume bounded segment curvature $\kappa_t \leq \kappa_{\max}$ and subspace approximation error $\|e_t\| \leq \varepsilon$ in each segment. Then piecewise low-rank reconstruction error on the sphere is*

bounded by a second-order curvature term plus the accumulated subspace error:

$$d_{\text{geo}}(h_t, \hat{h}_t) \leq C_1 \kappa_{\max} L_j^2 + C_2 \sum_{t \in S_j} \varepsilon_t, \quad (13)$$

where L_j is segment length in geodesic arc length.

Proof sketch. Approximate the trajectory within each segment by integrating the projected tangent increments. Geodesic integration on the sphere has second-order local truncation error in curvature. The projection residual contributes linearly through Grönwall-style accumulation along the segment. Summing the two terms yields the stated bound. $\square$

B. Paper 2: Priority Allocation Under Scarcity

Proposition 2. *Suppose decoder harm is monotone in a geometry-derived risk score $r(u)$ and memory cost is additive across units. Then for fixed budget B , allocating memory to higher-risk units before lower-risk units weakly dominates uniform allocation in expected decoder distortion.*

Proof sketch. Under monotone marginal harm, the cost-constrained optimization becomes a knapsack-style prioritization problem. Any reallocation that moves memory from a lower-risk unit to a higher-risk unit cannot increase expected harm. Iterating this exchange argument yields weak dominance over uniform allocation whenever the risk ordering is informative. $\square$

C. Paper 3: Sparse Codec Stability and Signal Dependence

Theorem 2. *Assume (i) the decoder is locally Lipschitz in a decoder-aware geometry metric, (ii) compressed segment representations preserve anchors and the latest support turn with bounded reconstruction error, and (iii) omitted residual information is bounded by a fixed exogenous segment-wise harm predictor. Then decoder drift under keep/compress/drop is bounded by the sum of per-segment codec distortion terms.*

Proof sketch. The decoder Lipschitz condition turns state reconstruction error into logit drift. The sparse codec decomposes reconstruction error into an anchor term, a support-memory term, and an evicted-residual term. Bounding each component separately and summing over segments yields a codec-level distortion bound. The bound holds for any scoring signal that correctly identifies high-risk turns; the empirical evidence now shows that the correct signal is geometry for constraint-critical tasks and query-conditioned geometry within a semantic shortlist for persona-continuity tasks. $\square$

Remark 1. *The signal-dependence is not a failure of generality. It reflects a mathematical fact about the different information structures of the two memory types. For constraint-critical turns, the user constraint and its assistant echo have nearly identical cosine similarity, so semantic scoring cannot distinguish them; geometric state displacement can. For persona-continuity turns, semantic proximity to the target query is the primary ranking signal; geometry refines but does not replace it.*

TABLE I
CURRENT MODEL REGISTRY USED ACROSS THE CHECKPOINT.

Preset	Hugging Face model	Current role
qwen25_05b	Qwen/Qwen2.5-0.5B-Instruct	Paper 1/Paper 2 local baseline
qwen25_15b	Qwen/Qwen2.5-1.5B-Instruct	Main fairness-controlled model
qwen25_3b	Qwen/Qwen2.5-3B-Instruct	First 3B Paper 2/Paper 3 validation
smollm2_17b	HuggingFaceTB/SmolLM2-1.7B-Instruct	Non-Qwen compact control family
llama32_3b	meta-llama/llama-3.2-3B-Instruct	Non-Qwen 3B public-benchmark runner

IV. EXPERIMENTAL DESIGN

A. Models

We evaluate compact and mid-scale open models suitable for local and Colab-based study. Table I summarizes the current preset registry.

B. Conversation Families and Evaluation Sets

Paper 1 uses audited short multi-family conversations for geometric characterization and changepoint analysis. **Paper 2** uses a harder long-context stress set with three families: long dependency, retrieval-heavy, and code conversation. The hard set is deliberately constructed so that late-turn answers depend on earlier facts or constraints. **Paper 3 hard-set experiments** reuse the same hard set. **Paper 3 Gate 1 and signal-comparison experiments** additionally use two public benchmarks:

- **MSC valid** (gonced8/multi-session-chat, 16 conversations used): multi-session dialogue with persona continuity, preference tracking, and topic flow. Tests semantic-memory recall, not constraint preservation.
- **LongMemEval-S cleaned** (xiaowu0162/longmemeval-cleaned): longer episodic-retrieval conversations with distributed relevance (500 conversations, 398–618 turns each; truncated to first 40 turns per conversation for tractable oracle computation). Used for Gate 1 LongMemEval oracle (6-conversation subset, 10,416 candidate rows) and Gate 1 refinement study.

C. Policies

- **uniform**: budget spread without geometry.
- **geometry**: direct geometry-based retention.
- **semantic**: cosine-similarity-to-target turn selection.
- **geometry_segment_actions**: segment keep/compress/evict bridge policy.
- **geometry_KCD**: sparse codec with geometric scoring.
- **semantic_KCD**: sparse codec with semantic scoring (new; Experiment 1).
- **budget_aware_semantic_KCD**: semantic-led codec with budget-aware support selection.
- **semantic_ambient_geometry_KCD**: semantic shortlist + ambient geometry refinement inside the codec.
- **semantic_query_conditioned_geometry_KCD**: semantic shortlist + query-projected geometric refinement (Gate 1 primary).

TABLE II
PAPER 1 MAIN QUANTITATIVE RESULTS. RANK95 REMAINS SMALL, SHUFFLED-TURN NULLS INCREASE RANK, AND GEOMETRY-LOGIT COUPLING REMAINS EXTREMELY HIGH.

Model	Rank95	95% CI	Shuffled Rank95	p_{H1}	$\text{Corr}(d_{\text{geo}}, \Delta\ell)$	p_{H3}
qwen25_05b	1.049	[1.000, 1.111]	1.486	0.0000	0.989	0.0000
qwen25_15b	1.167	[1.083, 1.271]	1.458	0.0008	0.994	0.0000
smollm2_17b	1.528	[1.500, 1.576]	1.799	0.0122	0.989	0.0000

TABLE III
PAPER 2 GEOMETRY-VS-UNIFORM RESULTS ON THE HARD STRESS SET. NEGATIVE DELTAS ARE IMPROVEMENTS OVER UNIFORM.

Model	$\Delta\text{Logit}@0.20$	p	$\Delta\text{Logit}@0.35$	p	$\Delta\text{NLL}@0.35$	p
qwen25_05b	-65.654	0.0000	-99.503	0.0000	-0.091	0.7222
qwen25_15b	-33.647	0.0015	-21.669	0.2637	-0.181	0.1343
smollm2_17b	-4.105	0.2878	-11.329	0.0200	-0.002	1.0000

D. Statistical Protocol

The repository computes bootstrap standard deviations and 95% confidence intervals for key means, along with paired sign-flip or paired-difference tests relative to uniform and pairwise policy comparisons where applicable. This manuscript reports only tracked results already present in versioned artifacts and the Drive-persisted Colab run outputs.

V. RESULTS

A. Paper 1: Conversation-State Geometry Is Compact and Decoder-Relevant

The core Paper 1 result is strong and stable. Conversation-state trajectories are extremely low-rank across all tested models, with mean rank95 between 1.049 and 1.528. Shuffling turn order substantially increases rank95, ruling out the null explanation that low-rank structure is merely an artifact of bounded dimensionality. Geometric distortion is almost perfectly aligned with output-space drift.

For qwen25_05b, mean rank95 is 1.049 [1.000, 1.111], shuffled null 1.486, geometry-logit correlation 0.989. For qwen25_15b: rank95 1.167, correlation 0.994. For smollm2_17b: rank95 1.528, correlation 0.989. These are not borderline effects. They define the scientific core of the program.

B. Paper 2: Geometry-Aware Control Beats Uniform Under Scarcity

On the hard stress set, geometry-aware control reduces decoder distortion relative to uniform retention. At budget 0.35, geometry improves mean logit drift by 99.50 on qwen25_05b and 21.67 on qwen25_15b.

The mechanism is concrete: at budget 0.35 on qwen25_05b, geometry retains more support user turns than uniform in 14/36 cases (uniform wins 5/36), and keeps the latest support turn while uniform drops it in 13/36 cases. The rescued turn types are exact code constraints, SQL ordering requirements, retrieval logistics, and compact fact bundles.

Paper 2 Checkpoint: Geometry-Aware Control

Geometry-guided retention remains the strongest decoder-fidelity policy on the hard stress set. The figure now emphasizes stable plots plus the support-turn mechanism instead of the earlier weak heatmap panel.

Paper 2 Checkpoint: Stable Geometry-Control Evidence




Fig. 2. Paper 2 checkpoint overview. Geometry-family policies consistently outperform uniform on the hard stress set. The mechanism is support-turn rescue, not generic recency preservation.

C. Paper 3 Hard-Set Codec Results

On the fairness-controlled qwen25_15b hard-set sweep, geometry_KCD is the strongest logit policy through the low/mid-budget region:

- 0.24: $\Delta_{\ell_2} = -38.272$, $R_{\ell} = 0.898$, $p = 0.0022$
- 0.28: $\Delta_{\ell_2} = -42.138$, $R_{\ell} = 0.883$, $p = 0.0008$
- 0.32: $\Delta_{\ell_2} = -56.165$, $R_{\ell} = 0.845$, $p = 0.0030$
- 0.38: $\Delta_{\ell_2} = -41.881$, $R_{\ell} = 0.871$, $p = 0.0213$

At budget 0.24, KCD uses slightly *less* realized token fraction than uniform while still beating it on decoder distortion. Behavior improvements are significant at budgets 0.32–0.50 ($\Delta_{\text{NLL}} = -0.39$ to -0.64 , all $p < 0.03$).

The 3B probe on qwen25_3b sharpens the regime split: KCD is strongest at 0.35 ($\Delta_{\ell_2} = -53.675$, $p = 0.0272$), while plain geometry retakes the lead at 0.50 ($\Delta_{\ell_2} = +24.357$, $p =$

TABLE IV
PAPER 3 FAIRNESS AND 3B CHECKPOINT. KCD IS STRONGEST IN THE SCARCITY REGIME; GEOMETRY RETAKES THE LEAD AT LOOSER BUDGETS.

Setting	Policy	Budget	Δ_{Logit}	p	Δ_{NLL}	p
qwen25_15b fairness	geometry_keep_compress_drop	0.24	-38.272	0.0022	-0.034	0.5282
qwen25_15b fairness	geometry_keep_compress_drop	0.32	-56.165	0.0030	-0.390	0.0270
qwen25_15b fairness	geometry_keep_compress_drop	0.50	7.050	0.6730	-0.638	0.0040
qwen25_3b probe	geometry_keep_compress_drop	0.35	-70.997	0.0070	—	—
qwen25_3b probe	geometry	0.50	-46.526	0.0333	—	—

0.0063, KCD vs. geometry). Under scarcity, the sparse codec wins; once budget loosens, direct geometry retention retakes the lead.

D. Experiment 1: The Geometry Signal Is Load-Bearing

Question. Is the geometry signal specifically necessary for constraint-critical memory tasks, or does the KCD codec

Paper 3 Checkpoint: Compression Is Real

Both geometry_segment_actions and geometry_keep_compress_drop are viable. Segment actions are stronger at some higher-budget logit settings; keep/compress/drop is the cleaner low-to-mid-budget mechanism policy and preserves memory-critical support structure more often.




Fig. 3. Paper 3 hard-set checkpoint. KCD is the strongest low/mid-budget sparse codec under fairness control; plain geometry becomes the stronger higher-budget fallback on the 3B probe.

structure do the work regardless of scoring signal?

Setup. We run uniform, semantic, geometry, geometry_KCD, and semantic_KCD on the full hard stress set (qwen25_15b, budgets 0.20/0.35/0.50). semantic_KCD is a new policy that uses the same KCD codec structure with semantic cosine-to-target scoring instead of geometric risk scoring.

Results. semantic_KCD is significantly worse than geometry_KCD at all budgets tested:

At 0.20 and 0.35, geometry_KCD beats semantic_KCD on both logit and answer quality with statistical significance. At 0.50, logit distortions tie but geometry_KCD still dominates on answer NLL ($\Delta = +0.628$, $p = 0.001$). Notably, semantic_KCD is also significantly worse than plain semantic at every budget on answer NLL, showing that the

TABLE V
EXPERIMENT 1: GEOMETRY_KCD VS. SEMANTIC_KCD ON THE HARD STRESS SET (QWEN25_15B). POSITIVE Δ MEANS SEMANTIC_KCD IS WORSE. GEOMETRY SIGNAL IS NECESSARY.

Budget	$\Delta\Delta_{\ell_2}$	p (row)	$\Delta\Delta_{\text{NLL}}$ (sKCD worse)
0.20	+37.8	0.028	+0.195 ($p = 0.001$)
0.35	+38.5	0.079	+0.424 ($p = 0.011$)
0.50	+0.1 (tie)	0.995	+0.628 ($p = 0.001$)

KCD codec structure with the wrong signal actively harms performance.

Mechanism. The memory-critical analysis at budget 0.35 reveals why:

The mechanism is mathematically interpretable. User con-

TABLE VI
SUPPORT-TURN RESCUE COMPARISON AT BUDGET 0.35:
GEOMETRY_KCD VS. SEMANTIC_KCD.

Metric	geometry_KCD	semantic_KCD
Rescues more support turns than uniform	17/36	7/36
Uniform beats policy on support turns	2/36	5/36
Keeps latest support turn vs. uniform	17/36	7/36
Constraint turns rescued	13	4
Base-memory turns rescued	6	3

TABLE VII
GATE 1 ORACLE: GEOMETRY RANKING IMPROVEMENT INSIDE THE
SEMANTIC SHORTLIST ON MSC VALID. DECISION THRESHOLD: Δ
KENDALL $\geq +0.03$ OR Δ TOP-5 RECALL $\geq +5$ PP.

Budget	Δ Kendall tau	Δ top-5 recall
0.20	+0.530	+0.025
0.35	+0.602	+0.038
0.50	+0.546	+0.025

straint turns such as “Remember these event details exactly: venue Maple Hall, start 6 PM, host Jordan, backup host Nia” and their immediate assistant echo “Stored: venue Maple Hall, 6 PM, Jordan, Nia” have nearly identical cosine similarity to any query about the event. Semantic scoring cannot rank the user constraint higher than its echo. Geometric state displacement can, because the user constraint creates a larger local state displacement than the acknowledgment response.

E. Experiment 2: Gate 1 — Geometry Adds Value Inside a Semantic Shortlist

Question. Inside a semantic shortlist on a persona-continuity benchmark, do geometric features add ranking value over semantic scoring alone? And does the geometry hybrid policy beat the strongest semantic-only codec on answer quality?

Setup. We ran the full Gate 1 decision surface on MSC valid (16 conversations, qwen25_15b, budgets 0.20/0.35/0.50). The oracle probe generated 23,928 candidate rows scoring each turn against oracle harm labels using three features: `semantic_score`, `query_geom_v2_risk` (query-conditioned geometry), and `combined_structural_score`. The refinement study compared four policies: `semantic`, `budget_aware_semantic_KCD`, `semantic_ambient_geometry_KCD`, and `semantic_query_conditioned_geometry_KCD`.

Gate 1 oracle result: PASS.

The Gate 1 threshold was Δ Kendall $\geq +0.03$. Every budget exceeds the threshold by a factor of 15–20 $\times$. This is not a marginal pass. Within the semantic shortlist, `query_geom_v2_risk` improves top-5 oracle recall from 0.825 (pure semantic) to 0.850–0.863 across budgets, with reduced regret on the oracle top-5 selection.

Gate 1 policy result.

At budget 0.35, `semantic_query_conditioned_geometry_KCD` beats `budget_aware_semantic_KCD` on both

TABLE VIII
GATE 1 REFINEMENT STUDY ON MSC VALID (16 CONVERSATIONS, QWEN25_15B). PRIMARY COMPARISON:
SEMANTIC_QUERY_CONDITIONED_GEOMETRY_KCD VS.
BUDGET_AWARE_SEMANTIC_KCD. NEGATIVE Δ MEANS THE FIRST
POLICY IS BETTER.

Budget	$\Delta\Delta_{\ell_2}$	p (conv)	$\Delta\Delta_{\text{NLL}}$	p (row)
0.20	−34.5	0.209	−0.019	0.429
0.35	−26.4	0.044	−0.057	0.020
0.50	+24.9	0.450	+0.008	0.316

TABLE IX
GATE 1 ORACLE ON LONGMEMEVAL-S CLEANED (6 CONVERSATIONS, 10,416 ROWS). SEMANTIC KENDALL IS FLAT (0.013); GEOMETRY ACHIEVES 0.50–0.53 INSIDE THE SHORTLIST. SAME THRESHOLD AS MSC: Δ KENDALL $\geq +0.03$.

Budget	Semantic τ	Geometry τ	$\Delta\tau$	Gate 1
0.20	0.013	0.527	+0.514	✓ (17 $\times$)
0.35	0.013	0.520	+0.507	✓ (17 $\times$)
0.50	0.013	0.502	+0.489	✓ (16 $\times$)

logit distortion ($p = 0.044$ conversation-level) and answer NLL ($p = 0.020$ row-level). The query-conditioning is specifically responsible: ambient geometry (`semantic_ambient_geometry_KCD`) does not produce the same gain—the query-projected geometry beats ambient geometry on logit at 0.35 by $\Delta = -49.8$ ($p = 0.004$), confirming that query-conditioned projection into the conversation’s tangent frame is the operative mechanism.

At budget 0.20 and 0.50 the pattern is directionally similar but does not reach significance, consistent with prior findings that query-conditioned geometry is most useful in the mid-budget scarcity regime.

F. Experiment 3: Gate 1 on LongMemEval-S — Cross-Benchmark Confirmation

Question. Does the Gate 1 geometry ranking advantage replicate on episodic-retrieval conversations, or is it specific to the persona-continuity structure of MSC?

Setup. Oracle probe on 6 LongMemEval-S cleaned conversations (qwen25_15b, conversations truncated to 40 turns each, budgets 0.20/0.35/0.50, 10,416 candidate rows). Refinement study comparing the same four policies as in Experiment 2.

Gate 1 oracle result: PASS at 17–23 $\times$ threshold.

Semantic scoring is essentially non-informative on episodic-retrieval turns (Kendall 0.013 $\approx$ random). Geometry achieves Kendall 0.50–0.53, a 40 $\times$ improvement in within-shortlist ranking ability. The overall (non-shortlist-restricted) Δ Kendall is 1.47–1.54. Gate 1 passes on LongMemEval even more decisively than on MSC.

Refinement policy result. At budget 0.50, `semantic_query_conditioned_geometry_KCD` beats `semantic_ambient_geometry_KCD` by $\Delta\Delta_{\ell_2} = -76.6$ ($p = 0.036$ row-level), confirming that query-conditioned projection matters on episodic retrieval.

TABLE X

EXPERIMENT 4: GEOMETRY_KCD VS. SEMANTIC_KCD ON THE HARD STRESS SET (QWEN25_3B). POSITIVE Δ MEANS SEMANTIC_KCD IS WORSE. GEOMETRY ADVANTAGE STRENGTHENS WITH MODEL SCALE.

Budget	$\Delta\Delta_{\ell_2}$	p (row)	$\Delta\Delta_{\text{NLL}}$	p (row)
0.20	+44.9	0.048	+0.494	0.008
0.35	+69.7	0.050	+0.853	0.005
0.50	-29.1	0.128	+1.126	0.0003

TABLE XI

3B ANSWER NLL RELATIVE TO UNIFORM AT EACH BUDGET. NEGATIVE = BETTER THAN UNIFORM. SEMANTIC_KCD IS THE ONLY POLICY THAT HARMS ANSWER QUALITY AT EVERY BUDGET.

Policy	$\Delta\text{NLL} @ 0.20$	$\Delta\text{NLL} @ 0.35$	$\Delta\text{NLL} @ 0.50$
geometry	-0.158	-0.036	-0.626
geometry_KCD	-0.134	-0.357	-0.775
semantic	-0.017	-0.147	-0.665
semantic_KCD	+0.360	+0.496	+0.350

Answer NLL is near-flat across policies due to the 40-turn truncation constraint (all policies retain essentially the same early-conversation context); logit distortion is the operative metric and geometry wins.

G. Experiment 4: 3B Model-Scale Validation

Question. Does the signal-comparison result from Experiment 1 hold at higher model capacity, or is it specific to 1.5B compact models?

Setup. Reruns the Experiment 1 signal comparison (uniform, semantic, geometry, geometry_KCD, semantic_KCD) on the full hard stress set with qwen25_3b (Qwen/Qwen2.5-3B-Instruct), 9 conversations, 540 evaluations, 270 behavior evaluations, budgets 0.20/0.35/0.50.

Results. The geometry advantage grows substantially at 3B.

At budgets 0.20 and 0.35, geometry_KCD beats semantic_KCD on both logit and answer NLL with statistical significance. At budget 0.50, the logit delta reverses (budget looseness makes both converge on logit), but answer NLL damage from semantic_KCD peaks at +1.13 nats ($p = 0.0003$) — the largest effect size in the entire program. The only policy that makes answer quality worse than uniform at every budget is semantic_KCD; all geometry policies consistently improve answer quality.

Mechanism. Support-turn rescue at budget 0.35 on 3B: geometry_KCD rescues more support turns than uniform in 15/36 cases (7/36 for semantic_KCD), an identical ratio to the 1.5B result. The mechanism is scale-invariant at the turn-selection level, while the consequence of getting it wrong grows: a 3B model’s richer representations create a larger downstream quality gap when wrong turns are evicted.

Implication. The reviewer concern that geometry-based scoring “only works on tiny models” is directly inverted. At 3B, semantic scoring becomes more harmful, not less, because the richer representational space makes the distinction between constraint turns and their echoes more pronounced.

TABLE XII

OBJECT-LEVEL VS. TURN-LEVEL COMPRESSION ON MSC VALID (QWEN25_15B). OBJECT POLICIES ARE SIGNIFICANTLY WORSE; TOKEN FRACTIONS REVEAL CHRONIC UNDER-RETENTION.

Policy	Logit L2 @ 0.35	NLL @ 0.35	Token frac @ 0.35
semantic	928.9	2.438	0.743
budget_aware_sem_KCD	931.2	2.487	0.736
sem_object_KCD	970.9	2.979	0.519
sem_object_harm_KCD	1023.0	3.251	0.431

Geometry’s ability to exploit that distinction scales with model capacity.

H. Open-Item Investigation: Learned Harm Predictor

A 2-head MLP (HarmPredictor) was trained from 23,928 MSC oracle ablation rows to predict logit and NLL damage when a turn is removed. The selected model (with_attention, 28 features including raw and sink-corrected attention weights) achieves row-level Kendall $\tau = 0.637$ on validation and 0.612 on held-out test — genuinely learning oracle-damage rank ordering. However, conversation-level Kendall is -1.0 on validation: the predictor assigns higher absolute harm scores to turns from conversations with lower oracle harm. When evaluated as a selection policy (semantic_harm_KCD), it loses to semantic_query_conditioned_geometry_KCD at budget 0.35 by $\Delta\Delta_{\ell_2} = +51.2$ ($p = 0.023$ row, $p = 0.004$ conversation-level).

The failure mode is cross-conversation calibration: point-wise MSE loss on harm scalars enforces within-conversation rank ordering but not between-conversation absolute scale. Geometric features avoid this problem because transported increments are computed on the same hypersphere regardless of which conversation they come from. The natural fix is a pairwise ranking loss (e.g., RankNet or ListMLE) with training pairs drawn across conversations. This is filed as future work.

I. Open-Item Investigation: Memory-Object Level Compression

Object-granularity KCD (semantic_object_KCD) groups turns into semantic objects (persona, event, constraint, update, generic bundles) using a marker-based classifier, then applies keep/compress/drop decisions at object granularity with type-specific preservation scores (constraint: 0.92, update: 0.89, event: 0.88, persona: 0.87, generic: 0.84). On MSC valid, the result is unambiguously worse than turn-level baselines:

The object codec retains only 43–52% of the target budget’s tokens, leaving large fractions of the memory budget unused. The failure is mechanical: marker-based boundary detection over-segments conversations into many small objects; after within-object compression each retains 1–2 turns; total retained mass falls well below budget. The concept (object-level granularity) is sound; the implementation needs model-based boundary detection and tighter per-object budget accounting. Filed as future work.

TABLE XIII
COMPLETED PAPER 3 REGIME MAP. THE WINNING SCORING SIGNAL DEPENDS ON THE MEMORY TYPE. THE KCD CODEC STRUCTURE WORKS ACROSS BOTH REGIMES.

Evaluation set	0.20	0.35	0.50	Reading
Hard stress / fairness / 3B	KCD	KCD	geometry	Support-turn-critical memory; sparse geometry codecs help most under scarcity.
LongMemEval public	semantic	semantic / segment	KCD	Broader episode memory; geometry-family policies remain useful but ranking changes.
MSC	semantic family	semantic	semantic	Persona and conversational continuity dominate; geometry-KCD is the wrong family.
LoCoMo (bounded)	semantic family	semantic family	semantic family	Mixed temporal-semantic memory; semantic-led codecs remain strongest in current probes.




Fig. 4. Completed Paper 3 winner map. Hard support-turn benchmarks require geometric scoring; semantic-memory benchmarks require semantic scoring with query-conditioned geometric refinement. The KCD codec structure is common to both.

J. The Completed Benchmark Regime Map

The four new experiments, together with the prior public-benchmark studies and two investigated open items, complete the Paper 3 regime map.

The key finding is that the separation is principled, not benchmark-specific noise:

- **Constraint-critical memory (hard stress set):** geometry_KCD is necessary. semantic_KCD is actively harmful (Experiment 1). Semantic scoring cannot distinguish user constraints from assistant echoes.
- **Persona-continuity memory (MSC valid):** semantic_query_conditioned_geometry_KCD is optimal at budget 0.35. Geometry adds real value inside the shortlist (Gate 1 oracle Δ Kendall +0.53–0.60). Ambient geometry alone does not help; query-projection is necessary.
- **Episodic retrieval (LongMemEval-S cleaned):** Gate 1 passes at 16–17 $\times$ threshold (Δ Kendall +0.49–0.51 inside shortlist; semantic Kendall 0.013 $\approx$ random). At budget 0.50, semantic_query_conditioned_geometry_KCD beats ambient geometry by $\Delta\Delta_{\ell_2} = -76.6$ ($p = 0.036$). Query-conditioned geometry is necessary; ambient geometry is not sufficient.
- **Mixed temporal-semantic (LoCoMo):** Semantic-led fam-

TABLE XIV
NEGATIVE-RESULT CHECKPOINTS. THESE TESTS SHARPEN THE CLAIM BOUNDARY OF THE GEOMETRY PROGRAM.

Check	Setting	Quantitative result	Reading
MSC support vs filler curvature	5 conv.	$\Delta\kappa = 0.122$, 3/5 positive	No robust within-topic curvature separation.
Regime atlas	208 segments	retrieval_heavy: 6 in regime 1, 6 in regime 3	After stabilization, geometry-only regimes still mix stress and casual segments.
State-update alignment	10 conv. synthetic	mean $A = 0.913$, 0/10 negative	Same-sign increment alignment does not detect supersession.
State/increment cross-term	10 conv. synthetic	update 0.971 vs control 0.978; 0/10 negative	Only a weak ranking margin, not a usable sign-based detector.

ilies are currently safer; geometry helps under tight budgets and on event chains.

K. Negative Results and Explicit Limits

The checkpoint is stronger because several plausible extensions failed. Table XIV records the main falsification results.

a) *MSC persona/filler curvature.*: In a manual five-conversation support-vs-filler check, mean support-filler curvature delta was 0.122 (3/5 positive, 2/5 negative). Curvature does not separate semantically important persona turns from nearby filler inside MSC.

b) *Geometry-only regime atlas.*: After fixing a genuine long-conversation curvature blow-up (raw curvature up to 6314 on near-stationary segments, stabilized to 26), the atlas still failed to separate retrieval-heavy hard-set segments from MSC and LoCoMo. Geometry alone does not recover a usable benchmark taxonomy.

c) *State-update supersession.*: Both same-sign alignment ($A(s, t) = \cos(u_s, u_t)$, mean 0.913, 0/10 negative) and a geometrically cleaner state-vs.-increment cross-term (update cross 0.971, control cross 0.978) fail as sign-based detectors on a synthetic explicit-update benchmark. The model encodes both "user states a fact" and "user updates a fact" as similarly oriented increments; representational resolution is insufficient for clean sign separation in compact models.

d) *Query-conditioned geometry as standalone selector.*: At LoCoMo budget 0.35, plain query-conditioned geometry collapses to $\Delta_{\ell_2} = +881.4$; the KCD wrapper recovers to +9.7 but does not win. Query-conditioned geometry is a refinement signal inside a semantic shortlist, not a standalone primary selector.

VI. DISCUSSION

A. What Is Now Scientifically Settled

- 1) **Geometry is real and decoder-relevant.** Paper 1 rules out the strongest null explanations and shows extreme geometry-decoder coupling across model families.
- 2) **Geometry helps under memory scarcity on constraint-critical tasks.** Paper 2 establishes a real mechanism: support-turn rescue with turn-level interpretability.

- 3) **The KCD sparse codec is a real and active compression structure.** It is budget-responsive, fairness-controlled, and mechanism-valid.
- 4) **The geometry signal is load-bearing for constraint-critical memory.** Experiment 1 shows `semantic_KCD` is actively harmful on the hard stress set; the codec structure alone does not do the work.
- 5) **Geometry adds significant value inside a semantic shortlist, across benchmark families.** Gate 1 passes on MSC (Δ Kendall +0.53–0.60, 15–20 $\times$ threshold) and LongMemEval (Δ Kendall +0.49–0.53, 16–17 $\times$ threshold). The result is benchmark-general, not benchmark-specific.
- 6) **The geometry advantage grows with model scale.** At 3B, `semantic_KCD` is catastrophically harmful ($\Delta\Delta_{\text{NLL}} = +1.13$, $p = 0.0003$) while geometry-based policies consistently improve answer quality. The support-turn rescue mechanism is scale-invariant (15/36 vs. 7/36 at both 1.5B and 3B) while the quality consequence of semantic failure grows with capacity.
- 7) **The signal-conditioned KCD framework is the right architecture.** One codec structure, two scoring regimes matched to memory type. The signal is a parameter, not a flaw.
- 8) **Several plausible geometry-only extensions fail in compact models.** Those failures define the boundary of the claim.
- 9) **Learned harm predictor and object-level compression are diagnosed future directions.** Both were attempted and produced informative negatives with specific diagnoses: cross-conversation calibration failure for the predictor (fix: pairwise ranking loss); over-segmentation and budget misestimation for object compression (fix: model-based boundary detection).

B. What Remains Open

- 1) **Learned harm predictor with cross-conversation calibration.** A pairwise ranking loss (RankNet or ListMLE) with training pairs drawn across conversations is the specific architectural change needed. The current pointwise predictor achieves row-level Kendall 0.64 but has inverted conversation-level calibration; that inversion causes policy-level failure.
- 2) **Model-based object boundary detection.** The marker-based object segmentation over-fragments conversations, causing chronic budget under-use (43–52% token retention vs. 74% target). A model-based classifier for semantic-object boundaries would resolve this; the object-level KCD structure is otherwise sound.
- 3) **Larger-scale systematic validation.** The 3B result is strong but rests on one model variant. The failed compact-model branches (state-update supersession, persona/filler separation) may become viable at 7B+. Systematic sweep across the open-weight model registry is warranted.
- 4) **Codec theorem with a learned predictor.** The Paper 3 stability theorem assumes a bounded exogenous harm

predictor. Extending it to a learned predictor with generalization error bounds would complete the theoretical story.

C. Limitations

- Mechanism and codec evidence is strongest on the custom hard stress set, which was designed to expose support-turn failures.
- The 3B result is strong but rests on one model variant (Qwen2.5-3B-Instruct); broader multi-model 3B replication is warranted.
- Gate 1 policy significance at budget 0.35 is real but rests on 16 MSC conversations and 6 LongMemEval conversations; broader replication is warranted.
- The LongMemEval Gate 1 uses conversations truncated to 40 turns for computational tractability. Full-length conversation oracle results remain outstanding.
- Boundary-localization evidence from Paper 1 is secondary and should not be oversold as exact regime recovery.
- The codec theorem in Section III still assumes a bounded exogenous harm predictor rather than a learned one.

VII. REPRODUCIBILITY AND REPOSITORY PACKAGING

This manuscript is paired with a versioned repository checkpoint. All three experiments are fully reproducible from committed scripts and notebook:

- **Experiment 1 (signal comparison):** `scripts/run_paper3_signal_comparison_hardset.sh`
- **Experiment 2 (Gate 1 MSC):** `scripts/run_paper3_gate1_real.sh`
- **Experiment 3 (Gate 1 LongMemEval):** `scripts/run_paper3_harm_oracle_probe.sh` + `scripts/run_paper3_gate1_refinement_probe.sh`
- **Experiment 4 (3B validation):** `scripts/run_paper3_signal_comparison_hardset.sh` with `MODEL=qwen25_3b`
- **Open items runner:** `notebooks/rt_paper3_open_items_runner.ipynb` (all four open items in sequence, Drive persistence, reconnect recovery, IndexError fix for truncated-conversation target sampling)
- **Signal + Gate 1 Colab runner:** `notebooks/rt_paper3_signal_comparison_gate1_runner.ipynb`
- **Benchmark sources:** MSC valid from `gonced8/multi-session-chat`; LongMemEval-S cleaned from `xiaowu0162/longmemeval-cleaned`; both normalized via `scripts/prepare_public_benchmark_jsonl.py`
- Frozen artifact bundles for Papers 1–3 under `artifacts/`
- Manuscript assets generated from tracked JSON under `manuscript/generated/`

The build entry point is:

```
python scripts/build_manuscript_assets.py
tectonic manuscript/paper_checkpoint.tex
```

VIII. CONCLUSION

The RT program set out to determine whether hidden-state geometry of evolving conversation trajectories is a useful signal for conversational memory compression. The answer is now complete and precise.

Paper 1 establishes the geometric substrate: conversation-state trajectories are extremely low-rank (mean rank95 1.05–1.53) and almost perfectly decoder-coupled (geometry–logit correlation 0.989–0.994) across compact model families. Paper 2 turns this into a real control law: geometry-aware retention beats uniform on constraint-critical tasks by preferentially rescuing exact support turns that downstream answers depend on (17/36 vs. 5/36 at budget 0.35). Paper 3 turns the control law into a sparse codec family and closes the geometry-versus-semantics gap with four experiments: geometry scoring is necessary for constraint-critical memory because semantic scoring cannot distinguish user constraints from their assistant echoes (17/36 vs. 7/36 support-turn rescue, $p = 0.028$); query-conditioned geometry adds substantial and significant value inside a semantic shortlist for persona-continuity memory (Δ Kendall +0.53–0.60 on MSC, $p = 0.020$ NLL at budget 0.35); that oracle ranking advantage replicates on LongMemEval-S episodic retrieval (Δ Kendall +0.49–0.53, 16–17 $\times$ threshold); and the geometry advantage grows with model scale, with `semantic_KCD` at 3B producing +1.13 nats of answer NLL damage ($p = 0.0003$) while geometry consistently improves answer quality. Two additional open items—a learned harm predictor and object-level compression—were attempted and yield precise diagnoses for future work rather than current contributions.

The final contribution is not a universal codec. It is the signal-conditioned KCD framework: one compression structure (keep/compress/drop), two scoring regimes matched to memory type (geometric for constraints, semantic+query-geometry for persona/episodic tasks), a 3B validation showing the advantage strengthens at scale, and a falsification record that defines where the signal distinction matters and where geometry-only approaches fail. The north star of the program—not one universal geometry codec, but a signal-conditioned sparse codec family for conversational memory whose advantage grows with model capacity—is reached.

APPENDIX A

STORAGE ACCOUNTING FOR SPARSE CODECS

Any honest codec comparison must count more than just retained raw tokens. For segment j , the storage model should include the anchor state, sparse support memory, and metadata:

$$\text{Storage} = \sum_j \left(\underbrace{\|a_j\|}_{\text{anchor}} + \underbrace{\|s_j\|}_{\text{sparse support}} + \underbrace{\|\text{meta}_j\|}_{\text{segment metadata}} \right). \quad (14)$$

Future work should refine this accounting into exact byte counts for latent summaries and decoder-compatible surrogates.

TABLE XV
APPENDIX BRIDGE FROM THE CHECKPOINT MANUSCRIPT TO THE FULL DERIVATION LEDGER.

Layer	Status	Core evidence	Deep-dive companion
Paper 1	Frozen	Low rank + geometry–decoder coupling	ledger
Paper 2	Stable	Geometry-aware control + support-turn rescue	experiment matrix
Paper 3	Mixed	Hard-set wins; benchmark-dependent transfer	experiment matrix
Extensions	Failed / weak	MSC curvature, atlas, state-update tests	negative matrix

APPENDIX B

RESEARCH LEDGER BRIDGE

This checkpoint is paired with a companion derivation ledger at `papers/rt_derivation_ledger.md` and a full program findings synthesis at `papers/rt_full_program_findings.md`. The ledger is the canonical deep-dive reference for the full math-to-code-to-artifact chain: every mathematical object, every benchmark family, every major experiment phase, and every explicit falsification branch. Machine-readable inventories at `papers/generated/rt_experiment_matrix.csv` and `papers/generated/rt_negative_result_matrix.csv`. Table XV gives the condensed bridge from the manuscript-level claims to the deeper ledger.

APPENDIX C

CHECKPOINT ARTIFACT MAP

The most important tracked bundles are:

- Paper 1 final: `artifacts/paper1/expanded_v8_final`
- Paper 1 regime atlas limits: `artifacts/paper1/regime_atlas_smoke_v4`
- Paper 2 hard stress: `artifacts/paper2/behavior_stress_v1`
- Paper 2 mechanism cases: `artifacts/paper2/behavior_stress_qwen_cases`
- Paper 3 fairness: `artifacts/paper3/paper3_batch_v1_fairness`
- Paper 3 3B probe: `artifacts/paper3/paper3_batch_v1_3b`
- Paper 3 MSC persona falsification: `artifacts/paper3/msc_persona_curvature_v1`
- Paper 3 state-update falsification: `artifacts/paper3/state_update_alignment_smoke_qwen05b`
- Paper 3 state/update cross control: `artifacts/paper3/state_update_cross_control_qwen05b`

- Experiment 1 signal comparison (Drive):
rt_results/paper3/studies/paper3_
signal_comparison_hardset_v1
- Gate 1 MSC oracle (Drive): rt_results/paper3/
harm_oracle/paper3_gate1_oracle_msc_
valid
- Gate 1 MSC refinement study (Drive): rt_results/
paper3/studies/paper3_gate1_refinement_
msc_valid
- Gate 1 LongMemEval oracle (Drive): rt_open_items_
results/paper3/harm_oracle/paper3_
gate1_oracle_longmemeval_s_cleaned_
subset
- Gate 1 LongMemEval refinement (Drive):
rt_open_items_results/paper3/studies/
paper3_gate1_refinement_longmemeval_s_
cleaned_subset
- 3B signal comparison (Drive): rt_open_items_
results/paper3/studies/paper3_3b_
signal_comparison_v1
- Learned harm predictor models (Drive): rt_open_
items_results/paper3/harm_predictor_
models/paper3_harm_predictor_msc_v1
- Harm predictor policy study (Drive): rt_open_items_
results/paper3/studies/paper3_harm_
predictor_msc_v1
- Object compression study (Drive): rt_open_items_
results/paper3/studies/paper3_semantic_
object_msc_v1

REFERENCES

- [1] D. Wu, H. Wang, W. Yu, Y. Zhang, K.-W. Chang, and D. Yu, “Long-MemEval: Benchmarking Chat Assistants on Long-Term Interactive Memory,” *arXiv preprint arXiv:2410.10813*, 2024.
- [2] A. Maharana, D.-H. Lee, S. Tulyakov, M. Bansal, F. Barbieri, and Y. Fang, “Evaluating Very Long-Term Conversational Memory of LLM Agents,” *arXiv preprint arXiv:2402.17753*, 2024.
- [3] J. Xu, A. Szlam, and J. Weston, “Beyond Goldfish Memory: Long-Term Open-Domain Conversation,” *arXiv preprint arXiv:2107.07567*, 2021.
- [4] P. Kompally, “RT Geometry Memory Repository Checkpoint,” GitHub repository, 2026. <https://github.com/SteveMama/rt-geometry-memory>